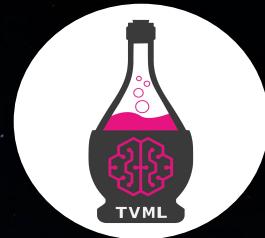


Introduction to Machine Learning with PyTorch

Cesare Montresor
cesare.montresor@gmail.com



TARALLUCCI, VINO E
MACHINE LEARNING
STUDY GROUPS

 School of AI
Verona

 iaml
ITALIAN ASSOCIATION FOR MACHINE LEARNING

Disclaimer

I'm a student, not an expert

...there will be some errors.

Why ?

- it's hype.
- Yields interesting results.
- It's here to stay.
- Solve problems otherwise impossible.
- Open new opportunities.
- Easy to implement O_o.

Idea

Give input and desired output and let the logic write itself (supervised learning).

Monet ↪ Photos



Monet → photo

Zebras ↪ Horses



zebra → horse

Summer ↪ Winter



summer → winter

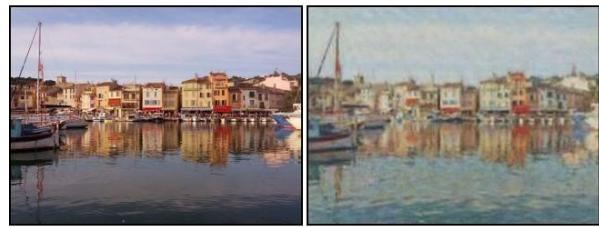


photo → Monet



horse → zebra



winter → summer



Photograph



Monet



Van Gogh

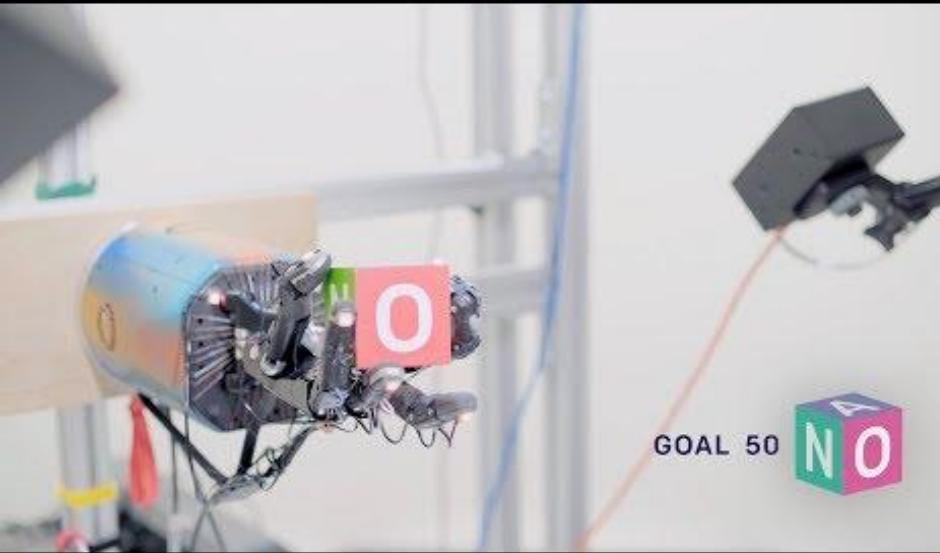


Cezanne



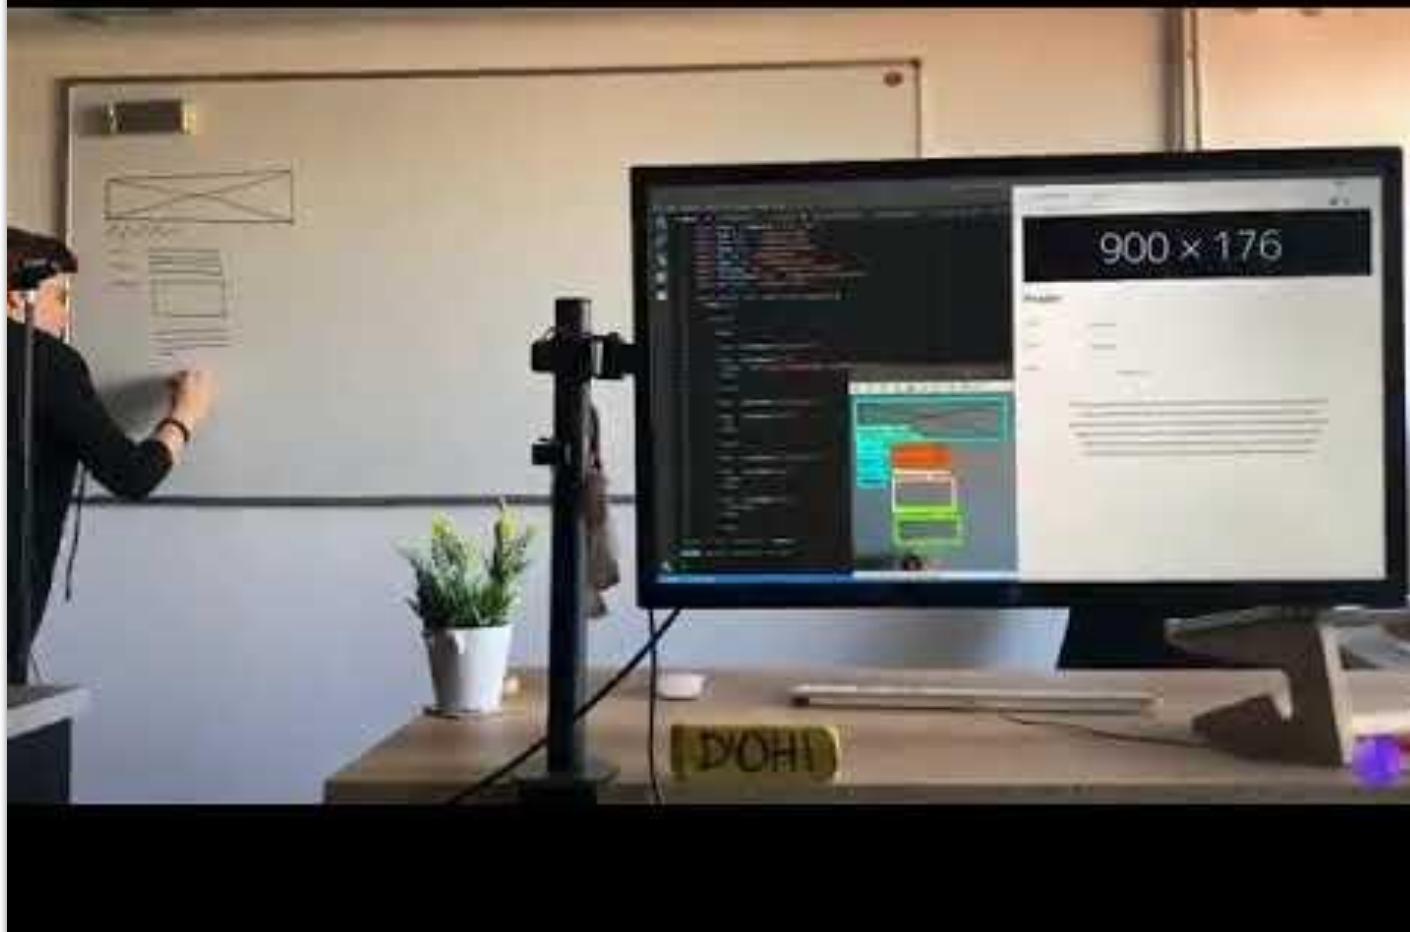
Ukiyo-e

OpenAI Dexterity



OpenAI Five





Steps

1. Processing the input.
2. Error compared to the desired output.
3. Correct slightly.

Repeat

1. Processing the input: a function

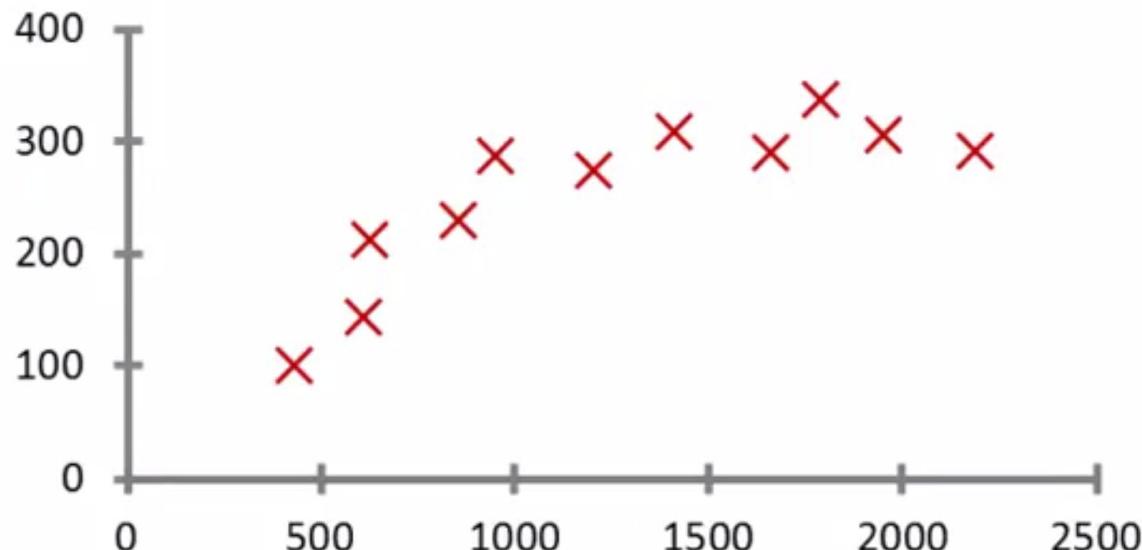
Input (x)	Output (y)
3	9
2	4
7	49
9	81

```
int y = myFunction(x)
```

$$y = f(x)$$
$$\text{myFunction}(x) == f(x) == x^2$$

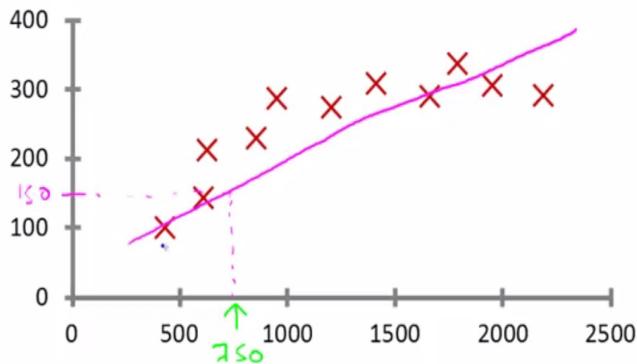
1. Example: real-estate price quotation

Y Price (€)	X Area (dm ²)
100.000	450
150.000	600
220.000	620
300.000	1000
...	...
290.000	2250



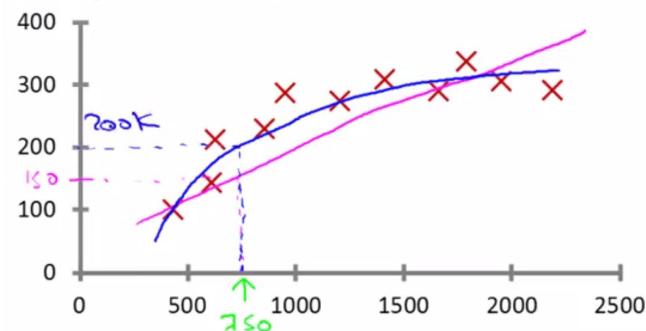
1. Approximation

Line



$$y = (x * a) + b$$

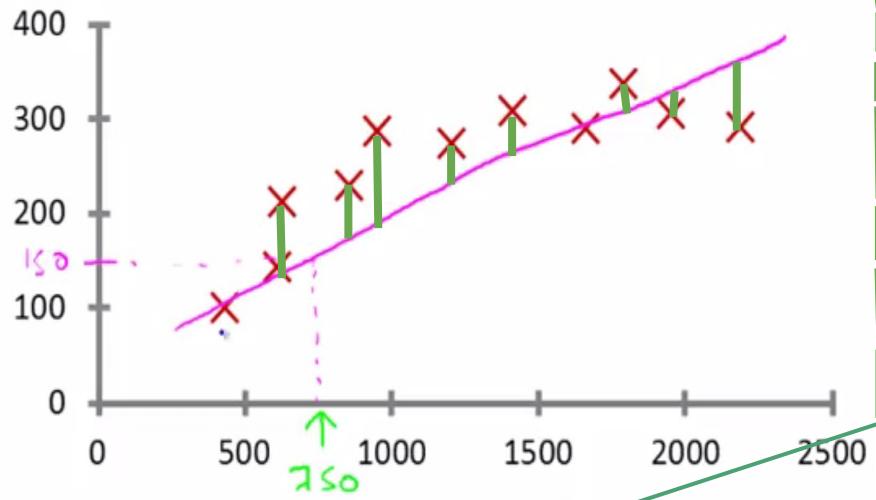
Parabola



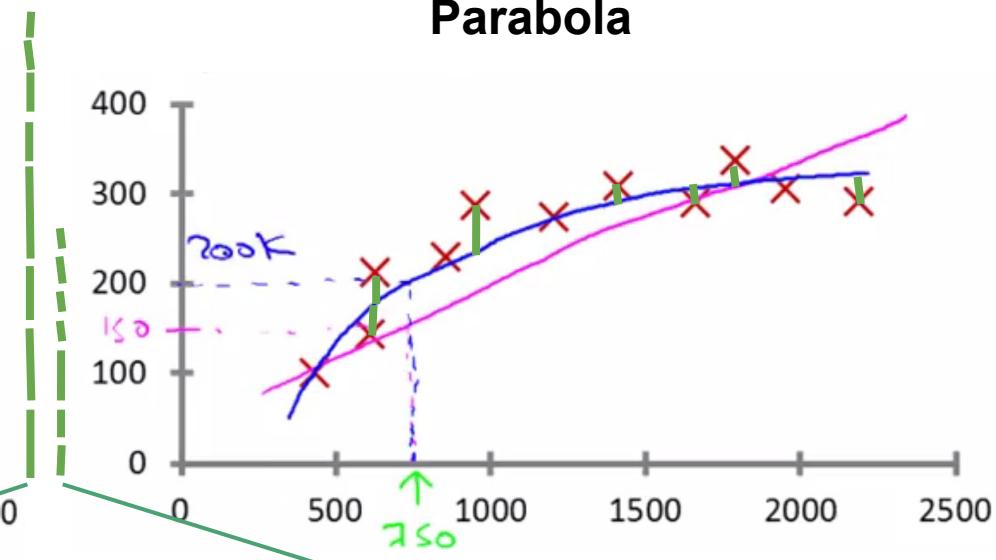
$$y = (x^2 * a) + (x * b) + c$$

2. Error (loss)

Line



Parabola



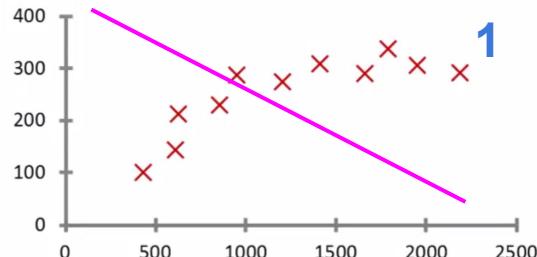
~500

~250

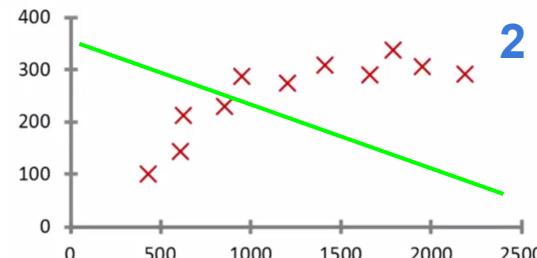
Nella pratica conviene usare MSE

$$\frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

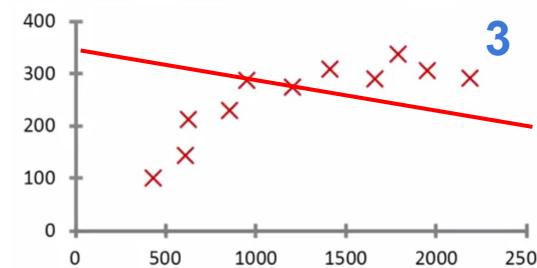
3. Correction (optimize)



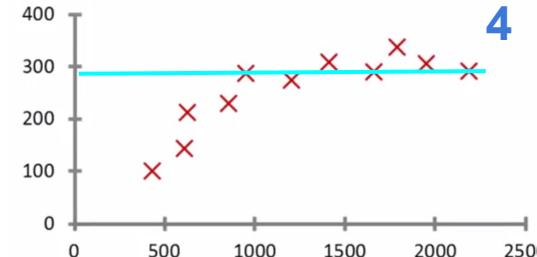
1



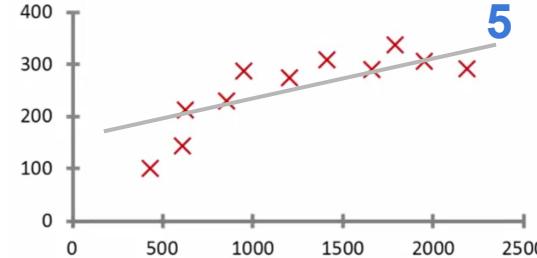
2



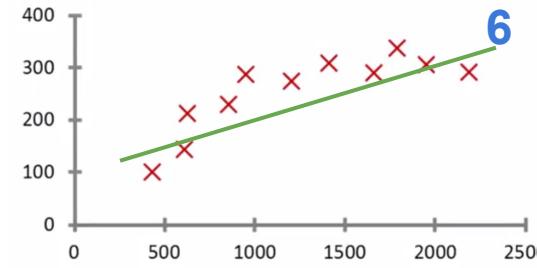
3



4

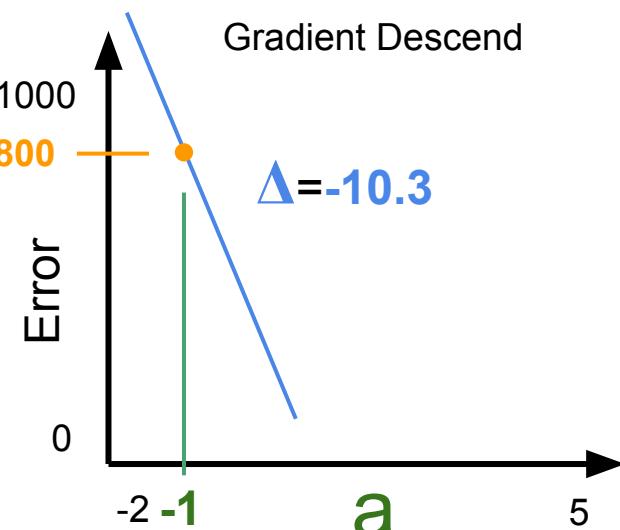


5



6

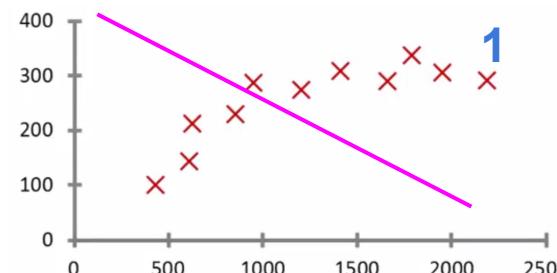
3. Correction (optimize)



Learning Rate (α) = 0.01

$$a = a - \Delta * \alpha$$

$$a = -1 - (-10.3) * 0.01 = \underline{-0.897}$$



$$y = (\mathbf{x} * \mathbf{a}) + b$$

$$\text{Errore} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

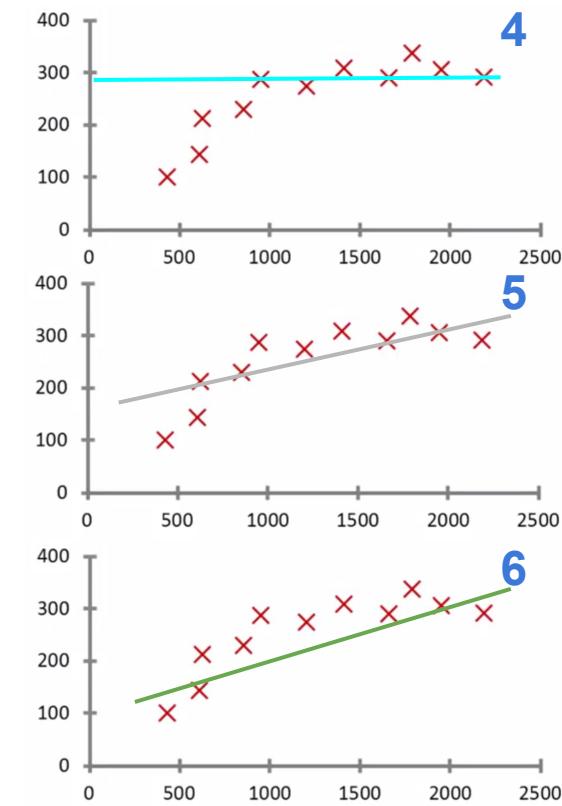
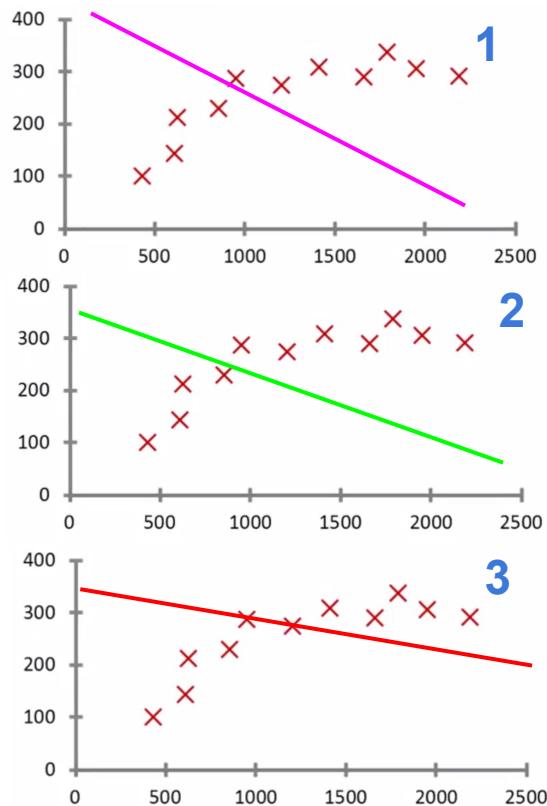
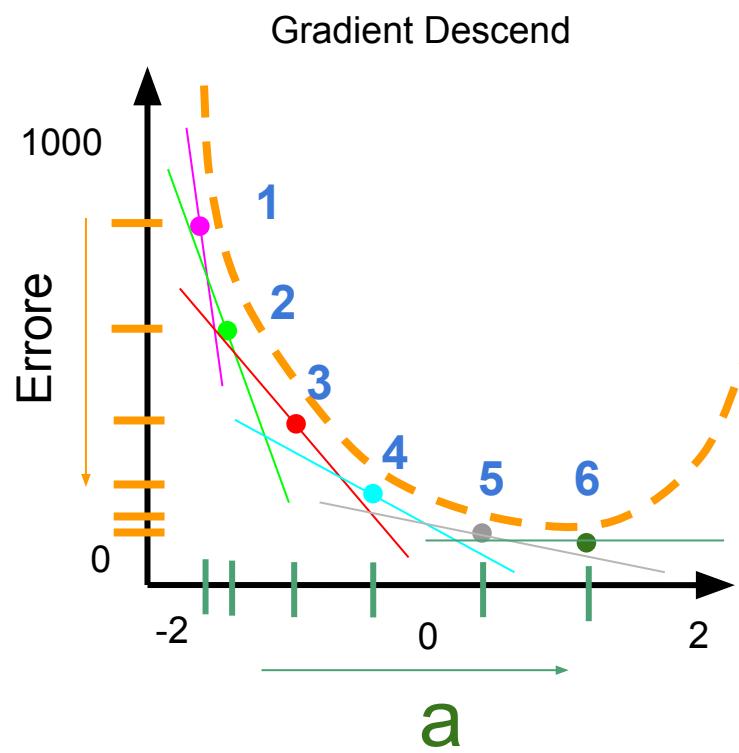
Error (loss):

$$((x_1 * a + b - \hat{y}_1)^2 + \dots + (x_n * a + b - \hat{y}_n)^2) / n$$

Derivative of the error with respect to a (Δ) "gradient"

$$(2 * x_1 * (x_1 * a + b - \hat{y}_1) + \dots + 2 * x_n * (x_n * a + b - \hat{y}_n)) / n$$

3. Correction (optimize)



```
from torch.utils.data import Dataset
import pandas as pd

class CsvDataset(Dataset):
    def __init__(self, file_path, transform=None):
        self.data = pd.read_csv(file_path)
        self.transform = transform

    def __getitem__(self, index):
        house_price = self.data.iloc[index, 0] # y
        house_info = self.data.iloc[index, 1:] # X
        if self.transform is not None:
            house_info = self.transform(house_info)
        return house_info, house_price

    def __len__(self):
        return len(self.data)
```

```
from torchvision import transforms
from torch.utils.data import DataLoader

my_dataset = CsvDataset('data.csv')

data_loader = DataLoader(
    my_dataset,
    batch_size = 8,
    shuffle = True
)

x,y = next(iter(data_loader))
```

```
import torch
import torch.nn.functional as F

x = torch.from_numpy( np.random.randn(10) )
y = torch.from_numpy( np.random.randn(10) )

a = torch.randn(1, requires_grad=True)
b = torch.randn(1, requires_grad=True)

y_pred = x*a+b

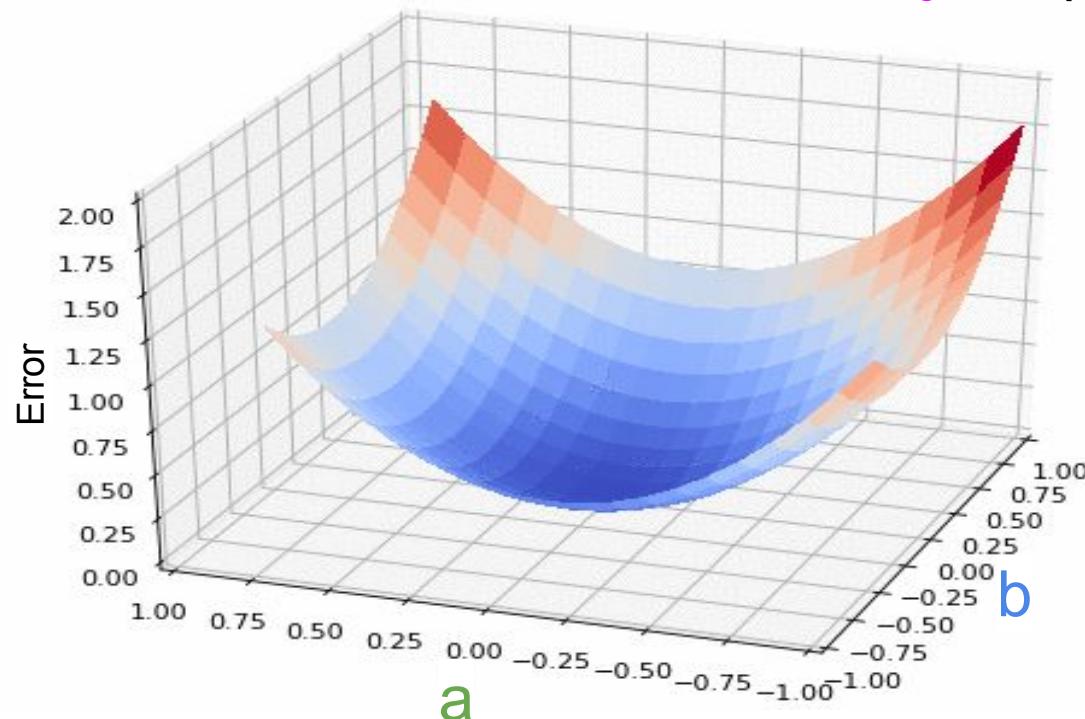
error = ((y_pred - y) ** 2).mean()
# Oppure, più semplicemente
error = F.mse_loss(y, y_pred)
```

```
learning_rate = 1e-3  
  
# x: house_info, y: house_price  
for x, y in train_loader:  
    y_pred = x*a+b  
    # calcolo dell'errore  
    error = ((y_pred - y) ** 2).mean()  
    # gradient usando la derivata dell'errore  
    error.backward()  
    # ottimizzo a e b  
    a.data -= learning_rate * a.grad.data  
    b.data -= learning_rate * b.grad.data  
    # reset gradient a 0  
    a.grad.data.zero_()  
    b.grad.data.zero_()
```

```
import torch.nn.functional as F  
from torch.optim import SGD  
from torch import nn  
  
learning_rate = 1e-3  
optimizer = SGD( [a, b], lr=learning_rate)  
model = nn.Linear(1,1)  
for x, y in train_loader:  
    y_pred = model(x)  
    # calcolo dell'errore  
    error = F.mse_loss(y_pred, y)  
    # gradient usando la derivata dell'errore  
    error.backward()  
    # ottimizzo a e b  
    optimizer.step()  
    # reset gradient a 0  
    optimizer.zero_grad()
```

For 2 parameters

$$y = (x * a) + b$$



More parameters

Price (y)	Year (x_1)	Area (x_2)	Var N (x_N)
300.000	1995	120	...
120.000	2001	83	...
49.000	1973	50	...
400.000	2011	200	...
...

```
int y = myFunction(x1,x2);  
y = f(x1,x2)
```

Artificial parameters:
 $x_3 = \sqrt{x_1^2 + x_2^2}$

$$y = f(x_1, x_2, \dots, x_N)$$

How many parameters ?



Size: ~200 Kb
Width: 1600 px
Height: 900 px
1 px = 255, 255, 255

$$1600 \times 900 \times 3 = \\ \mathbf{4.320.000}$$

(more than four millions)

↓
GPU

$$y = f(x_1, x_2, \dots, x_{4.320.000})$$

```
import torch

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

a = torch.Tensor((2,3), device=device)
model = torch.nn.Linear(10,10).to(device)
```

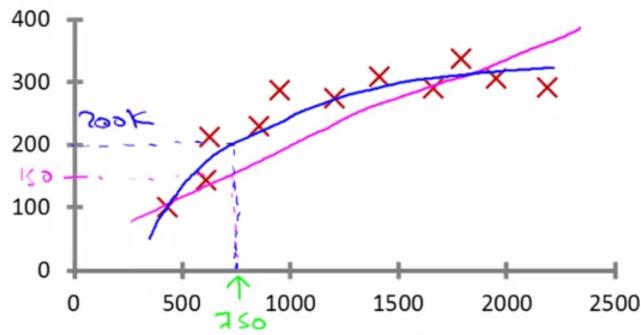
```
gpu_count = torch.cuda.device_count()

gpu0 = torch.device("cuda:0")
gpu1 = torch.device("cuda:1")
...
cpu0 = torch.device("cpu:0")
cpu1 = torch.device("cpu:1")
```

Categories of algorithms

- Regression / Classification
 - Supervised / Unsupervised
-

Regression

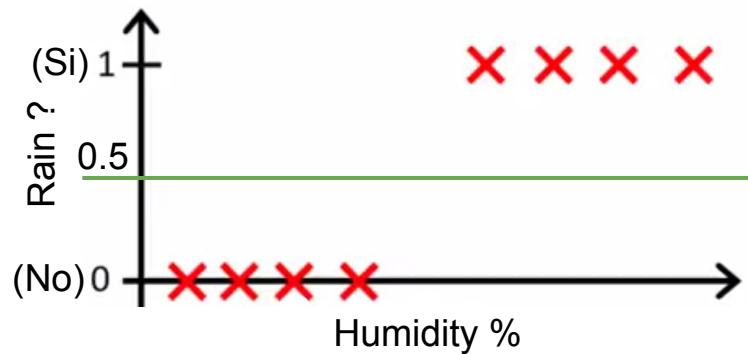


Obtain a continuous output.

Examples:

- Value of a good
- Probability of success.
- Control of an actuator.
- Age from a picture.

Classification

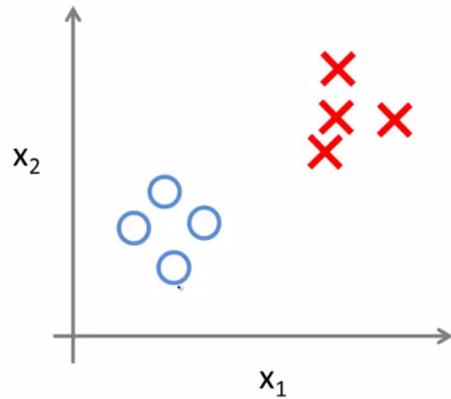


To determine a class/category

Examples:

- Kind of object, from picture
- Weather forecast.
- Gender from picture.
- Sentiment from text.

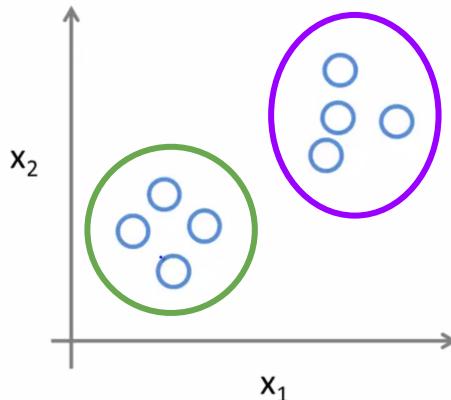
Supervised



Provided input and desired output.

- Linear Regression
- SVM
- Decision tree
- Neural Networks

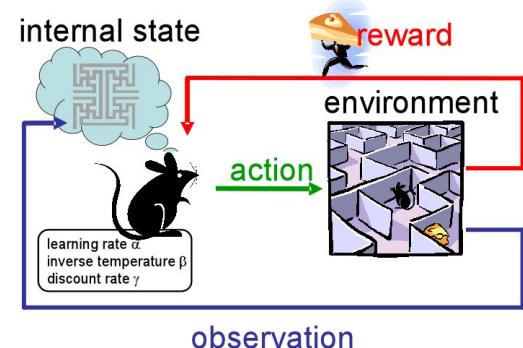
Unsupervised



Providing input and identify group/structure in data.

- K-means, dbscan
- PCA, SVD
- Anomaly detection
- Neural Networks

Reinforcement



Developing strategies based on actions and rewards.

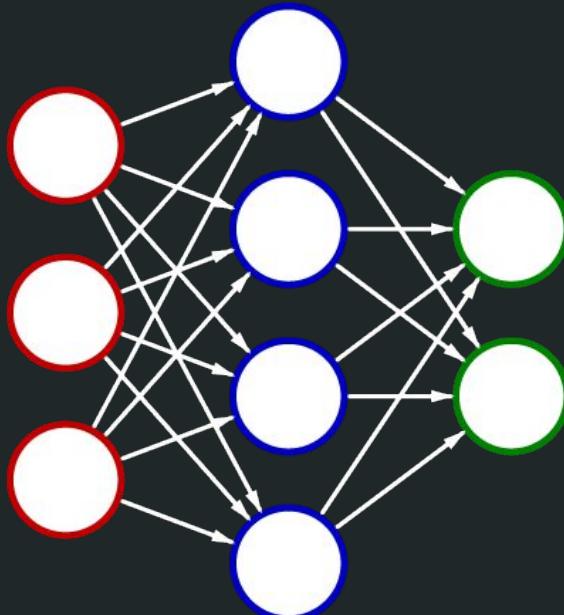
- Monte-Carlo method
- Q-Learning
- A3C, A2C
- Proximal Policy Optimization

Neural Networks

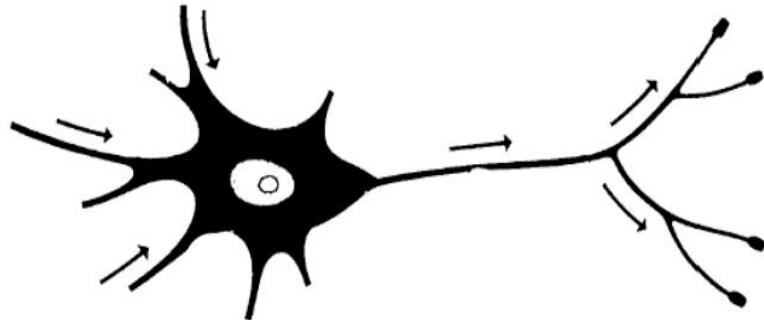
(Differentiable Universal Function Approximators)

(Differentiable Computational Graphs)

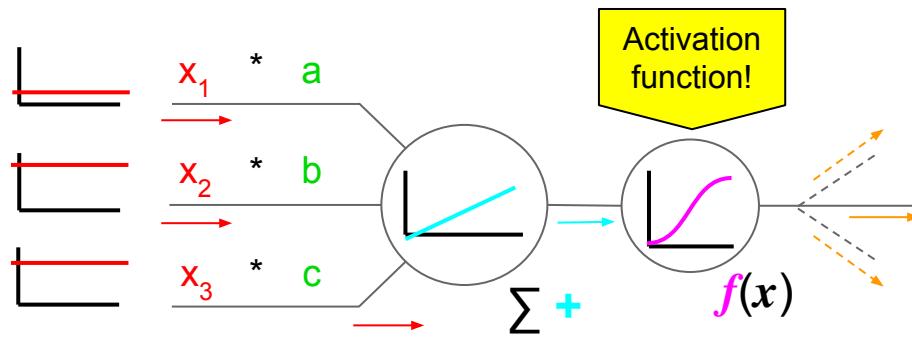
(≈deep learning)



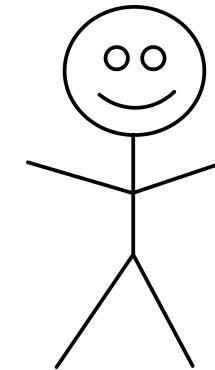
“Neuron”



=

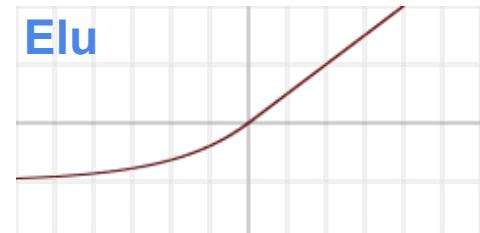
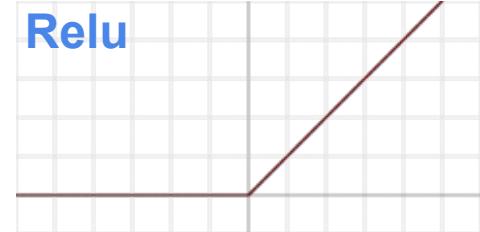
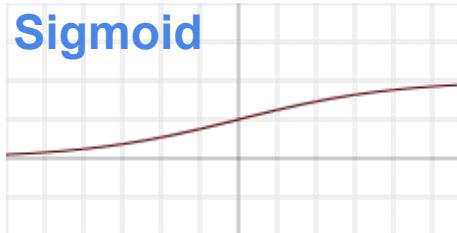


=



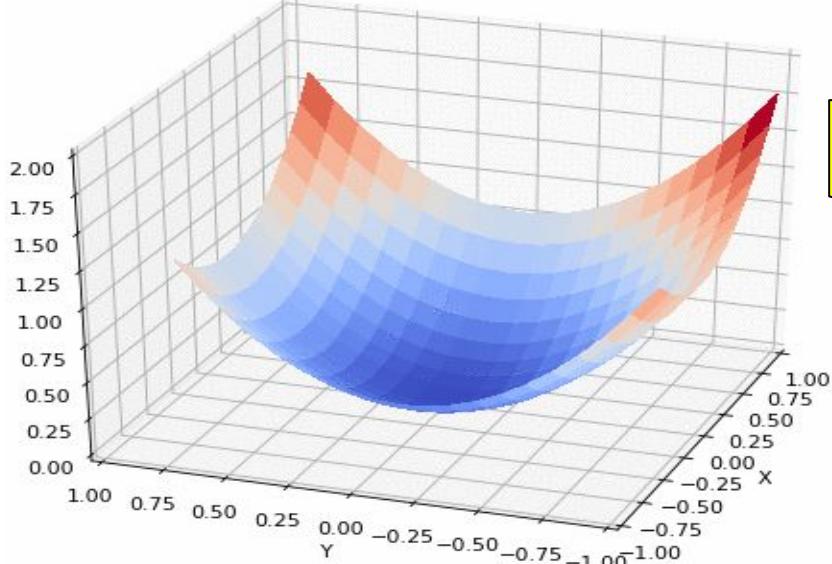
```
import torch.nn.functional as F  
  
a = torch.rand(4, requires_grad=True)  
  
F.relu(a)      # [0, ∞)  
F.sigmoid(a)   # (0,1)  
F.tanh(a)      # (-1,1)  
F.softmax(a)   # exp probability  
F.leaky_relu(a)  
F.elu(a)
```

Wikipedia Activation Function



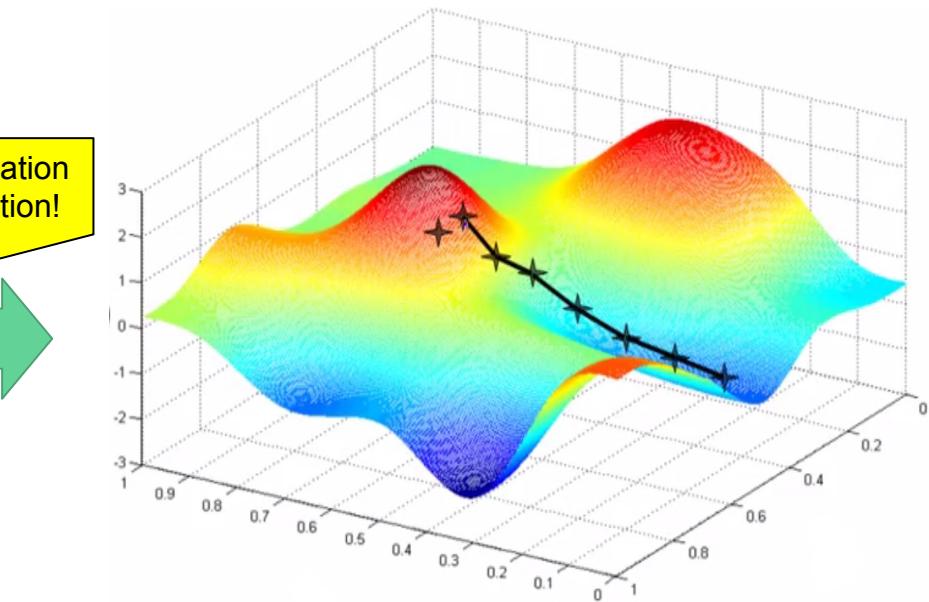
Optimizing non-linear functions

Improves the representation capabilities of the model, but makes non-deterministic the optimization process.



One global minima

Guaranteed to obtain the optimal solution



Many local minima

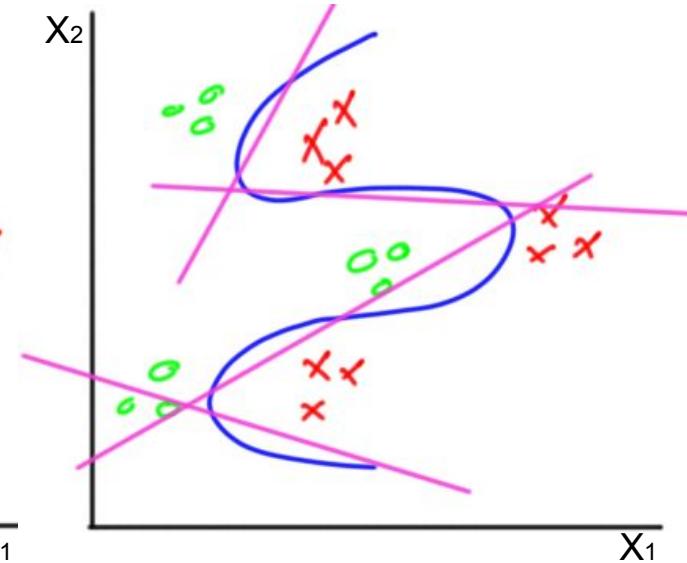
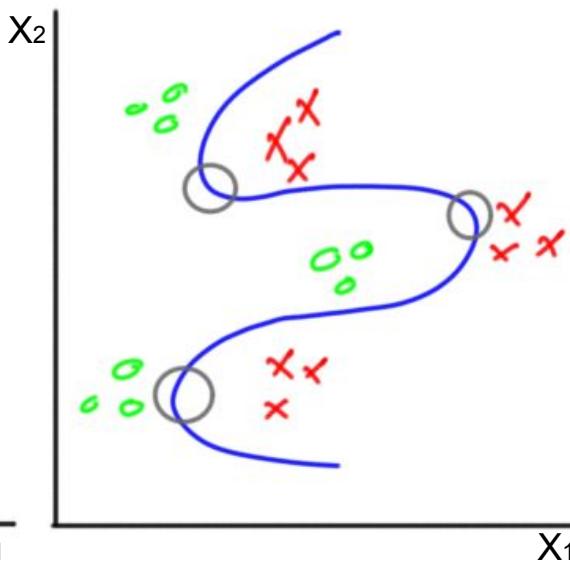
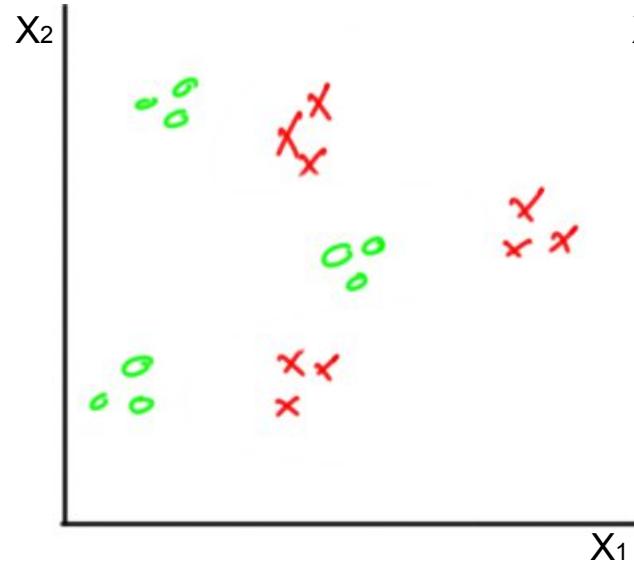
No guaranteed that any solution is the most optimal

Hierarchy of neurons

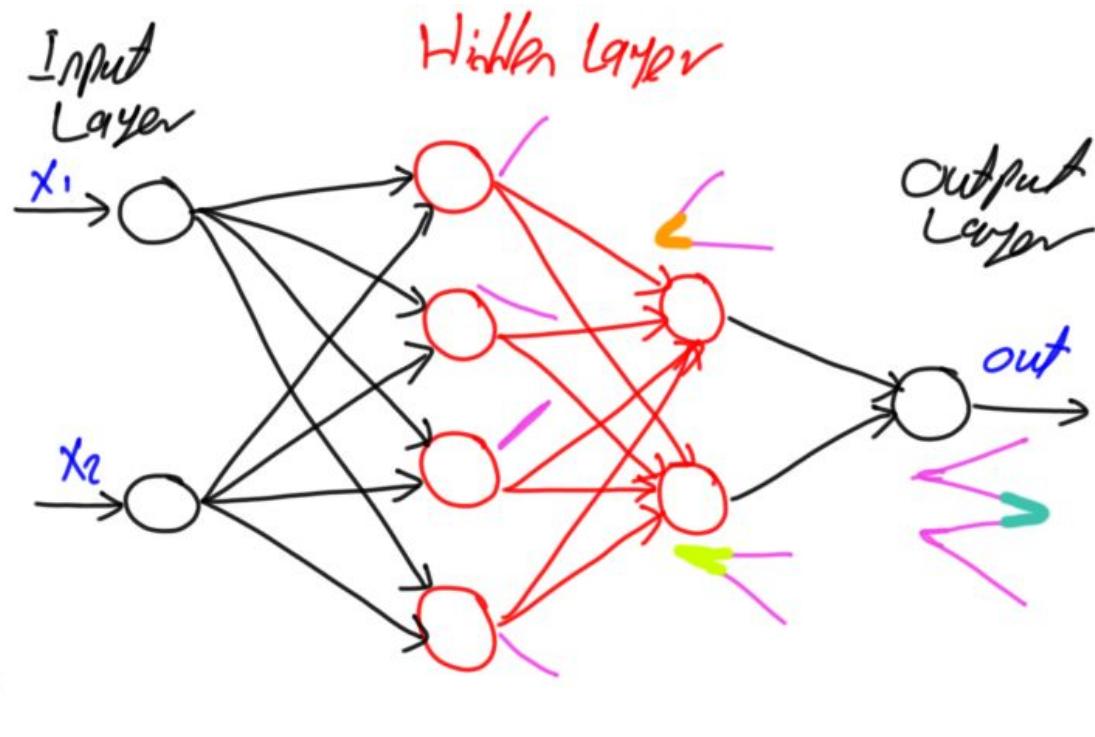
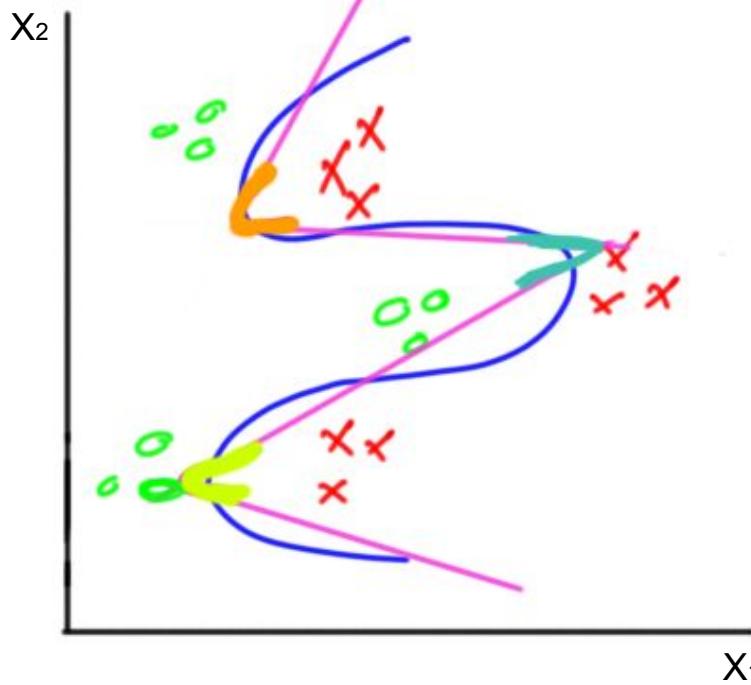
Problem: Identify issues in a engine.

X_1 = Vibrations

X_2 = Temperature



Hierarchy of neurons



```
from torch import nn
```

```
model = nn.Sequential(  
    nn.Linear(2, 4),  
    nn.ReLU(),  
    nn.Linear(4, 2),  
    nn.ReLU(),  
    nn.Linear(2, 1),  
    nn.Sigmoid()  
)
```

```
from torch import nn  
  
class NetworkFC(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.input = nn.Linear(2, 4)  
        self.hidden = nn.Linear(4, 2)  
        self.output = nn.Linear(2, 1)  
    def forward(self, x):  
        x1 = F.relu( self.input(x) )  
        x2 = F.relu( self.hidden(x1) )  
        x3 = F.sigmoid( self.output(x2) )  
        return x3  
  
model = NetworkFC()
```

Architectures

- Backfed Input Cell
- Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Different Memory Cell
- Kernel
- Convolution or Pool

A mostly complete chart of Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org

Perceptron (P)



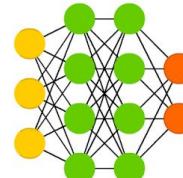
Feed Forward (FF)



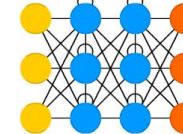
Radial Basis Network (RBF)



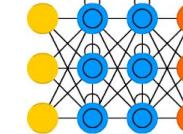
Deep Feed Forward (DFF)



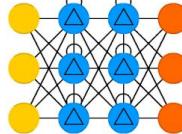
Recurrent Neural Network (RNN)



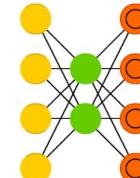
Long / Short Term Memory (LSTM)



Gated Recurrent Unit (GRU)



Auto Encoder (AE)



Variational AE (VAE)



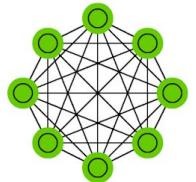
Denoising AE (DAE)



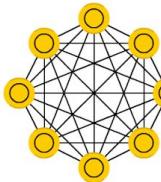
Sparse AE (SAE)



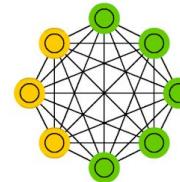
Markov Chain (MC)



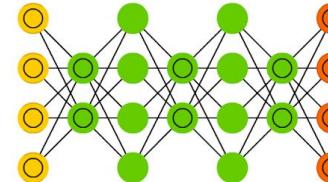
Hopfield Network (HN)



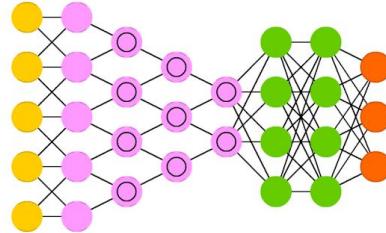
Boltzmann Machine (BM) / Restricted BM (RBM)



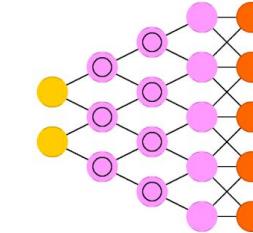
Deep Belief Network (DBN)



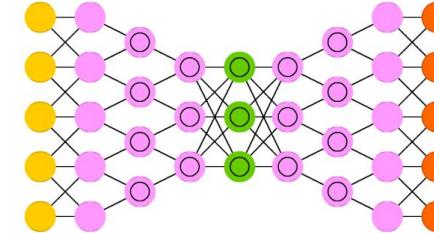
Deep Convolutional Network (DCN)



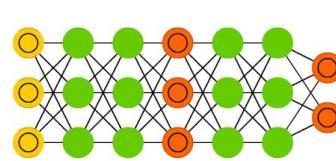
Deconvolutional Network (DN)



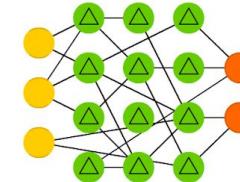
Deep Convolutional Inverse Graphics Network (DCIGN)



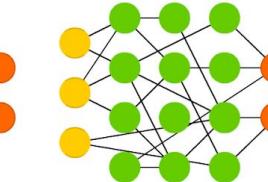
Generative Adversarial Network (GAN)



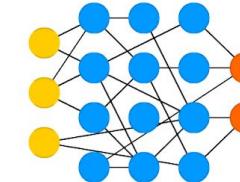
Liquid State Machine (LSM)



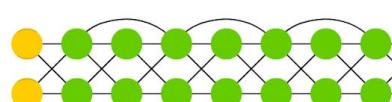
Extreme Learning Machine (ELM)



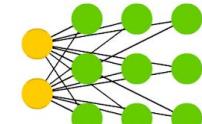
Echo State Network (ESN)



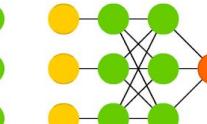
Deep Residual Network (DRN)



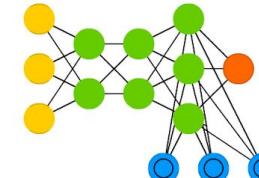
Kohonen Network (KN)



Support Vector Machine (SVM)

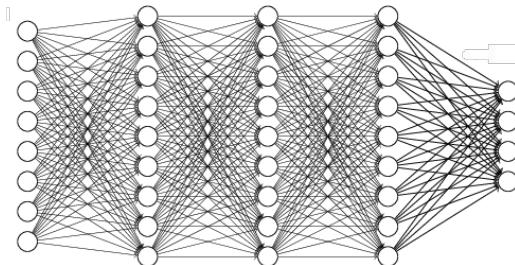


Neural Turing Machine (NTM)



NN

“Fully connected”

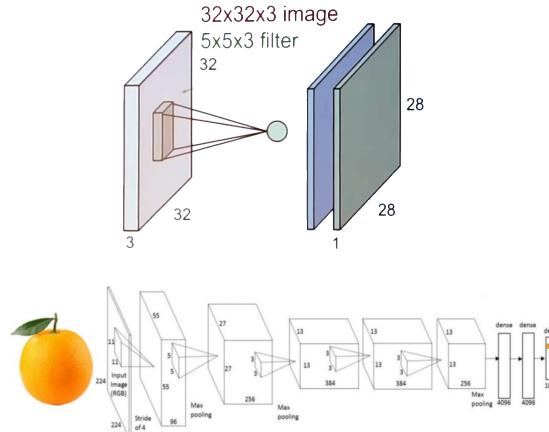


Every neuron is connected to every other neuron of the layers before and after

Also known as Linear, Dense, “fully connected” or “fc”.

CNN

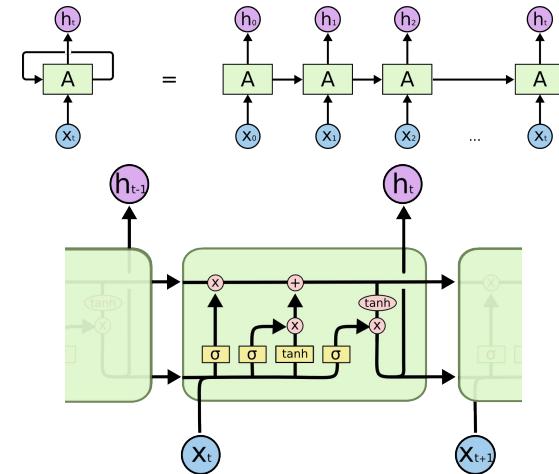
“Convolutional Neural Network”



Preserves the spatial information. Useful in parsing 2D/3D images, but also text.

RNN

“Recurrent Neural Networks”



Store an internal state while parsing sequential information. Useful in parsing time series, music, text, etc

```
from torch import nn
```

```
class torch.nn.Linear( in_features, out_features, bias=True)  
model = nn.Linear( 2, 4)
```

```
class torch.nn.Conv1d( in_channels, out_channels, kernel_size, ... )  
class torch.nn.Conv2d( in_channels, out_channels, kernel_size, ... )  
class torch.nn.Conv3d( in_channels, out_channels, kernel_size, ... )  
model = nn.Conv2d( 3, 20, 3)
```

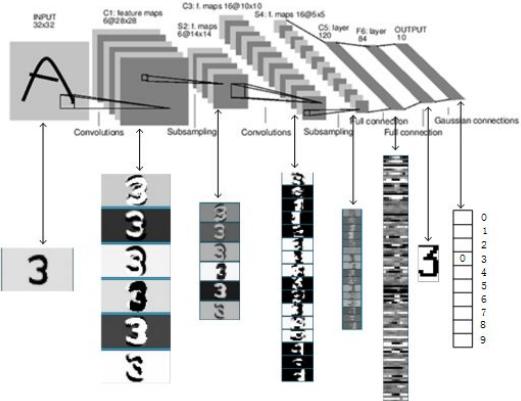
```
class torch.nn.RNNCell(input_size, hidden_size, ... )  
class torch.nn.LSTMCell(input_size, hidden_size, ... )  
class torch.nn.GRUCell(input_size, hidden_size, ... )  
class torch.nn.RNN(input_size, hidden_size, num_layers, ... )
```

```
model = nn.RNN(10, 20, 2)
```

LeNet

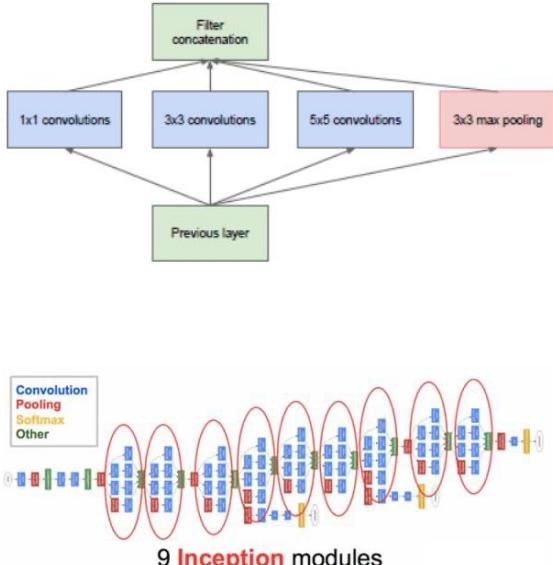
“Funzional!”

3 6 8 / 7 9 6 6 9 1
6 7 5 7 8 6 3 4 8 5
2 1 7 9 7 1 2 8 4 6
4 8 1 9 0 1 8 8 9 4
7 6 1 8 6 4 1 5 6 0
1 5 9 2 6 5 8 1 9 7
2 2 2 2 3 4 4 8 0
0 2 3 8 0 7 3 8 5 7
0 1 4 6 4 6 0 2 4 3
7 1 2 8 1 6 9 8 6 1



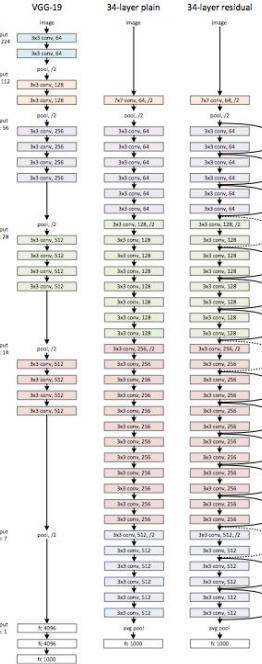
GoogLeNet

“Network nel network”



ResNet

“Up to 101!”



```
import torchvision.models as models

# AlexNet
model = models.alexnet()
# returns a model pre-trained on ImageNet
models.alexnet(pretrained=True)

# Inception v3 / GoogleNet
models.inception_v3( ... )

# SqueezeNet
models.squeezenet1_0( ... )
models.squeezenet1_1( ... )

# DenseNet
models.densenet121( ... )
models.densenet169( ... )
models.densenet161( ... )
models.densenet201( ... )

# VGG
models.vgg11( ... )
models.vgg11_bn( ... )
models.vgg13( ... )
models.vgg13_bn( ... )
models.vgg16( ... )
models.vgg16_bn( ... )
models.vgg19( ... )
models.vgg19_bn( ... )

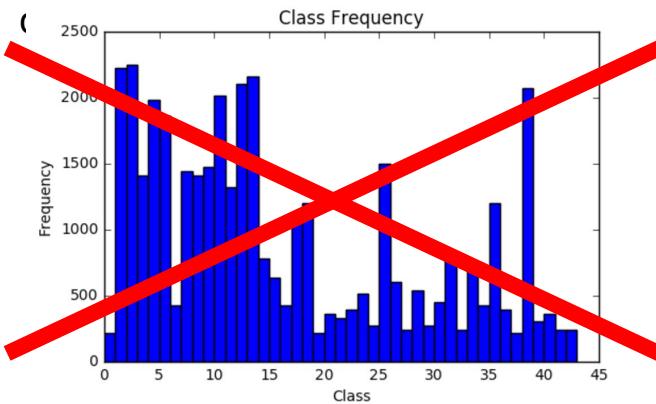
# ResNet
models.resnet18( ... )
models.resnet34( ... )
models.resnet50( ... )
models.resnet101( ... )
models.resnet152( ... )
```

... Dataset !!!

Often forgotten (1 slide ^_^), in favour of the architectures, the Dataset is of vital importance. It profoundly impact the quality of the learning and the final result.

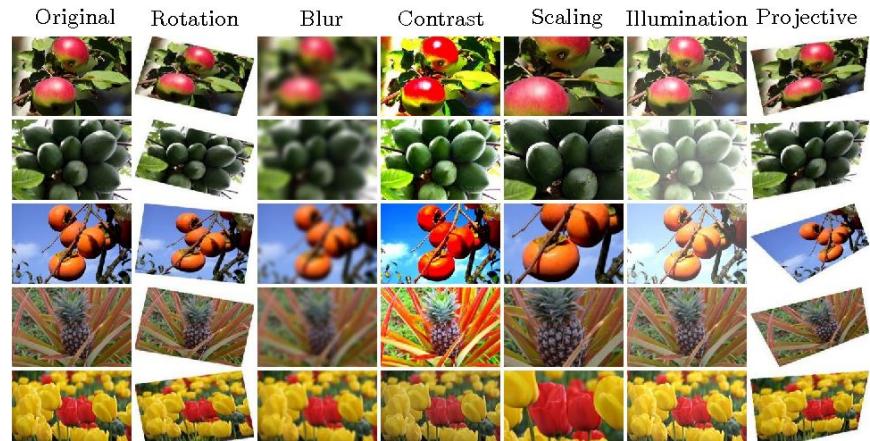
Wide, Varied and Balanced

- hundreds, thousents, millions ? it depends...
- Cats of every size, positions, colors, ages, dimentions, distance, etc...
- Equal numbers of elements. ($N_{\text{gatti}} \approx N$)



Augmenting

To extend the dataset with copies of the original data with added noise and distortions.



```
import torchvision.datasets as datasets
from torch.utils.data import DataLoader

mnist_ds = datasets.MNIST(root, train=True, download=True)
mnist_loader = DataLoader( mnist_ds, batch_size = 8, shuffle = True )
Images, labels = next(iter(data_loader))

# Other datasets
datasets.FashionMNIST( ... )
datasets.MNIST( ... )
datasets.CocoCaptions( ... )
datasets.LSUN( ... )
datasets.CIFAR10( ... )
datasets.STL10( ... )

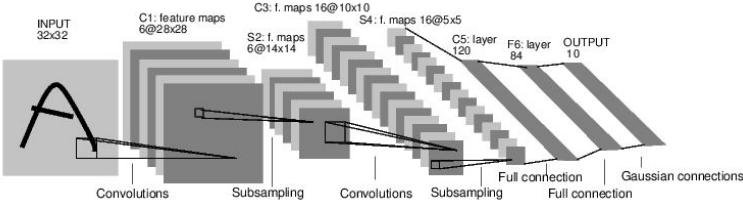
# Generic wrappers
datasets.ImageFolder(root, ... )
datasets.DatasetFolder(root, ... )
```

```
from torchvision import transforms
from torch.utils.data import DataLoader

transform_list = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize( mean=(0.5,), std=(0.5,) ),
    transforms.RandomHorizontalFlip(),
    transforms.RandomVerticalFlip(),
    transforms.RandomRotation( (-15, 15) ),
    transforms.RandomApply([
        transforms.ColorJitter(brightness=0.3)
    ])
])

mnist_ds = datasets.MNIST(root='.', train=True, transform=transform_list, download=True)
mnist_loader = DataLoader( mnist_ds, batch_size = 64, shuffle = True )
images, labels = next(iter(data_loader))
```

Questions ?



```
import torch
from torch import nn
from torch import optim
import torch.nn.functional as F
from torchvision import datasets, transforms

trainset = datasets.MNIST('.', download=True, transform=transforms.ToTensor())
trainloader = torch.utils.data.DataLoader(trainset, batch_size=256, shuffle=True)

model = nn.Sequential(
    nn.Conv2d(1, 20, kernel_size=5), nn.ReLU(),
    nn.Conv2d(20, 40, kernel_size=3, stride=2), nn.ReLU(),
    nn.Conv2d(40, 80, kernel_size=3, stride=2), nn.ReLU(),
    nn.Conv2d(80, 160, kernel_size=3, stride=2), nn.ReLU(),
    nn.AdaptiveMaxPool2d(1),
    nn.Linear(160, 10),
    nn.Softmax(dim=1)
)

optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

for epoch in range(10):
    for images, labels in trainloader:
        optimizer.zero_grad()
        output = model.forward(images)
        loss = F.mse_loss(output, labels)
        loss.backward()
        optimizer.step()

        loss = loss.item()
        print(f"Loss: {loss}", end="")
```

About

Cesare Montresor

- cesare.montresor@gmail.com
- <https://twitter.com/CesareMontresor>
- <https://linkedin.com/in/cmontresor>



School of AI

- <https://www.theschool.ai/>



School of AI
Verona

Tarallucci, Vino e Machine Learning

- <https://www.facebook.com/groups/TarallucciVinoMachineLearningVerona/>
- <https://www.meetup.com/Tarallucci-Vino-Machine-Learning/>



TARALLUCCI, VINO E
MACHINE LEARNING
GRUPPI DI STUDIO

Italian Association for Machine Learning

- <https://www.iaml.it/>
- <https://www.facebook.com/machinelearningitalia/>



Credits

images:

- <https://www.coursera.org/learn/machine-learning>
- <http://cs231n.stanford.edu>
- <https://github.com/junyanz/CycleGAN>
- <http://colah.github.io/posts/2015-08-Understanding-LSTMs>
- <https://www.slideshare.net/xavigiro/deep-convnets-for-global-recognition>
- <https://www.quora.com/How-does-the-Inception-module-work-in-GoogLeNet-deep-architecture>
- <https://www.quora.com/How-does-deep-residual-learning-work>
- <https://www.google.it/imghp>
- <http://www.asimovinstitute.org/neural-network-zoo>
- <https://www.linkedin.com/pulse/beginners-ask-how-many-hidden-layersneurons-use-artificial-ahmed-gad/>